

OrynticLabs Private Limited

Company & Product Overview Document

Version 1.0 | Confidential | For Internal and Partner Use

Who We Are

OrynticLabs is a full-spectrum technology company built for the era of intelligent software. We design, build, and deliver custom technology solutions across the complete software lifecycle — from early-stage product thinking to enterprise-scale deployment.

We are not a generalist IT vendor. We are a focused engineering and product studio that combines deep technical expertise with design intelligence and business understanding. Every solution we build is grounded in one belief: technology should solve real problems, move fast, and last long.

OrynticLabs operates across three delivery models — Software as a Service (SaaS), Platform as a Service (PaaS), and custom project-based engagements — giving clients the flexibility to work with us in whatever way fits their business best.

What We Do

1. Web Development

We build web applications, platforms, and digital products that perform at scale. Whether it is a consumer-facing product that needs to handle millions of users or an internal enterprise tool that needs to be reliable and fast, we approach every web build with the same standard: clean architecture, strong performance, and design that users actually enjoy using.

We work across the full web stack — frontend, backend, APIs, databases, and infrastructure — and we have deep experience with modern frameworks including Next.js, React, Node.js, Python (FastAPI, Django), and TypeScript.

We also work with leading content management and commerce platforms including Shopify, WordPress, Payload CMS, and headless CMS architectures — giving clients the ability to manage their digital presence independently, without depending on a developer for every update.

2. Product Development

We work with founders, product leaders, and enterprise teams to take ideas from concept to working product. This is not just about writing code. It includes product strategy, user research, information architecture, technical scoping, sprint planning, and iterative delivery — the full product development cycle.

We build MVPs for startups that need to validate quickly, and we build mature, scalable products for companies that have already validated and need to grow. We adapt to where you are in the journey.

3. Software Development

We develop custom software systems that are purpose-built for your business requirements. Off-the-shelf software solves generic problems. We build software that solves your specific problem — with the right architecture for your scale, the right integrations for your workflow, and the right flexibility for how your business will evolve.

This includes CRM systems, ERP platforms, internal tools, automation systems, workflow engines, and complex multi-service software architectures.

4. Mobile Development

We build native and cross-platform mobile applications for iOS and Android. Our mobile work covers consumer applications, enterprise mobility solutions, and hybrid products that work seamlessly across web and mobile.

We use React Native for cross-platform delivery and Swift / Kotlin for native builds where performance demands it. Every mobile product we deliver is designed for the device first — not adapted from a web layout.

5. Artificial Intelligence and Machine Learning

This is where OrynticLabs goes beyond what most technology companies offer. We build custom AI and ML systems — not wrappers around existing models, but intelligent systems trained on your data, integrated into your workflows, and designed to do real work.

Machine Learning (ML): Predictive models, classification systems, recommendation engines, anomaly detection, and customer intelligence — built on your historical data to generate decisions and insights specific to your business.

Deep Learning (DL): Neural network architectures for complex pattern recognition, natural language understanding, computer vision, and sequence modelling. We use deep learning where the problem genuinely requires it — and we are honest when a simpler approach will do the job better.

Applied AI: Custom AI agents, intelligent automation, large language model integrations, RAG (Retrieval-Augmented Generation) systems, and conversational AI — built for deployment in real production environments.

6. Data and Analytics

We help organisations make sense of their data. This covers data engineering (building the pipelines and infrastructure that collect, store, and process your data), data analytics (turning raw data into structured insight), and data visualisation (making that insight accessible to the people who need to act on it).

We build data warehouses, real-time analytics dashboards, business intelligence platforms, and event-tracking systems — and we connect all of it to the AI layer when the business is ready for it.

7. Cloud Computing and Infrastructure

We design and manage cloud infrastructure on AWS, GCP, and Azure. Our infrastructure work covers architecture design, cloud migration, containerisation with Docker and Kubernetes, CI/CD pipeline setup, security configuration, and ongoing infrastructure management.

We build infrastructure that is secure, observable, and as automated as possible — because infrastructure that requires constant manual attention is infrastructure that will eventually fail at the worst moment.

8. UI/UX Design

Design at OrynticLabs is not decoration. It is the layer that determines whether a product is used or abandoned. We bring a modern, premium design philosophy to every product we touch — combining clarity of structure, sophistication of aesthetic, and precision of interaction detail.

We deliver wireframes, interactive prototypes, complete design systems, and production-ready UI assets. Our design work is built to be implemented accurately — every spacing value, every component, every state is specified so that what gets built matches what was designed.

Industries We Serve

OrynticLabs has built products and platforms across a broad range of industries. Our approach in each is the same: understand the domain deeply before proposing a solution, and build something that fits how

that industry actually works.

Industry	What We Build
Fintech	Payment platforms, lending systems, portfolio management, fraud detection, KYC/AML automation, financial analytics
Healthcare and Biotech	Patient management systems, clinical data platforms, diagnostic AI, telemedicine infrastructure, biomedical data pipelines
EdTech	Learning management systems, adaptive learning platforms, student analytics, assessment engines, content delivery
Supply Chain and Logistics	Shipment tracking, warehouse management, route optimisation, supplier portals, inventory intelligence
Manufacturing	Production monitoring, quality control systems, predictive maintenance, factory floor digitisation
Agriculture	Crop intelligence, precision farming tools, supply chain visibility, market linkage platforms
Retail and E-commerce	Custom storefronts, inventory management, personalisation engines, omnichannel commerce, loyalty platforms
Energy and Utilities	Grid monitoring, energy consumption analytics, predictive maintenance, sustainability reporting
Real Estate	Property listing platforms, valuation tools, tenant management, digital agreements, broker portals
Automobile	Fleet management, vehicle telematics, dealer management systems, EV infrastructure software
Hospitality and Food Service	Reservation systems, POS integration, kitchen management, guest experience platforms, delivery operations

Technology Stack

We are technology-agnostic by principle. We choose the right tool for the problem — not the one that is trending or the one that is easiest to sell. What follows is a complete and honest picture of the technologies we work with, what each one is, and why it matters.

Frontend Technologies

The frontend is what users experience directly. We build frontend systems that are fast, accessible, and maintainable — not just visually impressive.

Next.js A React-based framework developed by Vercel. Next.js is our primary choice for production web applications because it supports server-side rendering (SSR), static site generation (SSG), incremental static regeneration (ISR), and the React Server Components model — giving us precise control over how pages are rendered and how fast they load. It is SEO-friendly, scalable, and integrates seamlessly with any backend.

React The industry-standard JavaScript library for building user interfaces. React's component model allows us to build complex, interactive UIs from small, reusable pieces. It is the foundation of our frontend work across web and mobile, and its ecosystem is the most mature in the industry.

TypeScript A strongly-typed superset of JavaScript. TypeScript catches errors at compile time rather than runtime, makes code easier to refactor, and significantly improves the developer experience in large codebases. We use TypeScript by default on every project — it is not optional for us.

Vue.js A progressive JavaScript framework used in cases where React is not the right fit. Vue is approachable, lightweight, and well-suited for projects where the team has existing Vue expertise or where the incremental adoption model is valuable.

Tailwind CSS A utility-first CSS framework that allows us to build custom designs without writing custom CSS for every element. Tailwind keeps styles consistent, reduces CSS bloat, and makes design changes faster to implement. We combine Tailwind with custom design tokens to maintain brand consistency across every component.

HTML5 and CSS3 The foundational standards of the web. We do not take shortcuts here — our HTML is semantic, our CSS is structured, and our markup is built for accessibility from the ground up.

Framer Motion and GSAP Animation libraries for React and vanilla JavaScript respectively. We use these for high-quality, performance-conscious animations and transitions — the detail that separates a premium product experience from a standard one.

Redux and Zustand State management solutions for complex frontend applications. Redux is used for large-scale applications with complex state logic. Zustand is our preferred choice for simpler, more maintainable state management in modern React applications.

React Query (TanStack Query) A server state management library that handles data fetching, caching, synchronisation, and background updates in React applications. It eliminates most of the boilerplate associated with API integration and significantly improves perceived application performance.

Backend Technologies

The backend is where the real logic of a product lives. We build backend systems that are reliable, secure, and designed to scale.

Node.js A JavaScript runtime built on Chrome's V8 engine. Node.js is well-suited for API servers, real-time applications, and microservices architectures. Its non-blocking I/O model handles high concurrency efficiently, and the shared JavaScript ecosystem between frontend and backend allows for faster development across the full stack.

Express.js The most widely-used web framework for Node.js. Express provides a minimal, unopinionated foundation for building HTTP servers and REST APIs. We use it for projects where flexibility and ecosystem breadth are priorities.

Python — FastAPI Our preferred Python framework for high-performance API development. FastAPI is built on modern Python standards including async/await and type hints. It generates automatic API documentation, provides excellent performance, and integrates naturally with the Python data science and AI ecosystem — making it the right choice for any backend that connects to ML or data systems.

Python — Django A high-level Python web framework built for rapid development and clean, pragmatic design. Django includes an ORM, authentication, admin interface, and many other features out of the box — making it well-suited for data-heavy applications and internal tools that need to be built and shipped quickly.

Go (Golang) A statically typed, compiled language designed at Google for high-performance systems. We use Go for backend services where raw performance, low latency, and efficient resource usage are critical — microservices that handle very high request volumes, real-time data processing pipelines, and CLI tooling.

Java and Spring Boot Java remains the dominant language in enterprise software. We use Spring Boot for enterprise-grade backend systems where the client's environment, existing team, or integration requirements make Java the right choice. Spring Boot provides a production-ready application framework with robust support for security, transactions, and enterprise integration patterns.

GraphQL A query language for APIs that allows clients to request exactly the data they need — no more, no less. We use GraphQL where the data access patterns are complex, the client-side flexibility requirement is high, or where multiple different clients (web, mobile, third-party) need to consume the same API in different ways.

REST API design Our default API design pattern for most projects. We follow RESTful principles rigorously — proper resource modelling, correct use of HTTP methods and status codes, versioning, rate limiting, and comprehensive documentation.

WebSockets and Socket.io For applications that require real-time bidirectional communication — chat systems, live dashboards, collaborative tools, notifications, and multiplayer features. Socket.io provides an abstraction over the WebSocket protocol with fallback support and room-based messaging built in.

Mobile Development

React Native A framework for building native mobile applications using React. React Native compiles to native UI components — not a web view — which means the performance and feel of the application is genuinely native, not a mobile-adapted web experience. One codebase runs on both iOS and Android, which significantly reduces development time and cost without sacrificing quality.

Expo A platform and toolchain built on top of React Native. Expo simplifies the development, testing, and deployment workflow for React Native applications. It provides a managed build service, over-the-air update capability, and access to native device APIs without writing native code — making it the right choice for most projects.

Swift Apple's native language for iOS and macOS development. We use Swift for iOS projects where performance, platform integration, or specific native APIs demand a fully native implementation. Swift is fast, safe by design, and gives us access to the latest Apple platform capabilities on day one.

Kotlin The modern language for Android development, fully supported by Google and the Android ecosystem. We use Kotlin for native Android development where the project requirements exceed what a cross-platform solution can offer.

Flutter Google's UI toolkit for building natively compiled applications for mobile, web, and desktop from a single codebase. Flutter uses the Dart language and its own rendering engine, which delivers consistent UI across platforms. We use Flutter for projects where pixel-perfect cross-platform UI consistency is a priority.

Artificial Intelligence and Machine Learning

Python (AI and ML ecosystem) Python is the lingua franca of AI and ML. Its ecosystem of libraries, frameworks, and community tooling is unmatched. Every AI or ML project we deliver is built in Python unless there is a specific reason to use another language.

TensorFlow Google's open-source machine learning framework. TensorFlow is well-suited for production ML deployments, particularly when models need to run at scale, be served via API, or be deployed to edge devices. We use TensorFlow for deep learning workloads and model serving.

PyTorch Facebook's open-source deep learning framework. PyTorch is the research community's preferred framework because of its dynamic computation graph and intuitive design. We use PyTorch for model development, experimentation, and production deployment — particularly for computer vision, natural language processing, and custom architecture work.

scikit-learn The standard Python library for classical machine learning. scikit-learn provides clean, consistent implementations of classification, regression, clustering, dimensionality reduction, and model evaluation tools. We use it for projects where deep learning is not required and a well-understood ML model is the right solution.

Hugging Face Transformers The industry-standard library for working with pre-trained language models. Hugging Face provides access to thousands of pre-trained models for tasks including text classification, named entity recognition, translation, summarisation, and question answering — and makes it straightforward to fine-tune these models on domain-specific data.

LangChain and LangGraph Frameworks for building applications with large language models. LangChain provides abstractions for chains, agents, memory, and tool use — the building blocks of LLM-powered applications. LangGraph extends this with graph-based agent orchestration for complex, multi-step reasoning workflows.

OpenAI API (GPT-4, GPT-4o) Access to OpenAI's frontier language models. We use the OpenAI API for conversational AI, content generation, code assistance, structured data extraction, and function calling — integrated into applications where the language model needs to take actions, not just generate text.

Anthropic Claude API Access to Anthropic's Claude family of models. Claude is particularly strong for tasks requiring long context windows, careful reasoning, and safe, predictable behaviour. We use Claude for document analysis, complex reasoning tasks, and applications where model reliability and safety are primary concerns.

Open-source language models (Llama, Mistral, Qwen) Self-hosted language models for use cases where data cannot leave the client's infrastructure, where cost at scale is a concern, or where a specific domain-tuned model performs better than a general-purpose frontier model. We deploy these models on client cloud infrastructure with full control over the compute, security, and access.

RAG — Retrieval-Augmented Generation An architecture that connects a language model to a knowledge base — allowing it to answer questions accurately based on specific documents, databases, or real-time data, rather than relying solely on what the model learned during training. We build RAG systems using vector databases (Pinecone, pgvector, Weaviate) combined with embedding models and retrieval pipelines.

Computer Vision — OpenCV, YOLO, Detectron2 For projects that require image and video understanding — object detection, image classification, optical character recognition, facial recognition, quality inspection, and more. We use OpenCV for image processing, YOLO for real-time object detection, and Detectron2 for complex scene understanding tasks.

Natural Language Processing Beyond large language models, we apply classical and neural NLP techniques for text preprocessing, sentiment analysis, entity extraction, topic modelling, and text classification — particularly for languages and domains where general-purpose LLMs underperform.

MLflow and Weights and Biases Model lifecycle management and experiment tracking platforms. We use these to track training experiments, version models, compare performance across runs, and manage the deployment of trained models to production — ensuring reproducibility and accountability in every ML project.

Celery A distributed task queue for Python. We use Celery to run ML inference jobs, model training tasks, and data processing pipelines asynchronously — ensuring that heavy AI computation does not block the main application thread.

Data and Analytics

Apache Spark A distributed computing engine for large-scale data processing. We use Spark for workloads that exceed what a single machine can handle — batch processing of large datasets, ETL pipelines, and distributed ML training at scale.

Apache Kafka A distributed event streaming platform. We use Kafka for real-time data pipelines, event-driven architectures, and systems that need to reliably process high-throughput streams of data — log ingestion, user event tracking, financial transaction streams.

dbt (data build tool) A transformation framework that brings software engineering best practices — version control, testing, documentation, modularity — to SQL-based data transformation. We use dbt to build reliable, maintainable analytics pipelines in data warehouses.

Google BigQuery Google's fully-managed serverless data warehouse. BigQuery is our preferred choice for analytics at scale — it can query terabytes of data in seconds, requires no infrastructure management, and integrates natively with Google's data ecosystem including Looker Studio and Vertex AI.

AWS Redshift Amazon's cloud data warehouse. We use Redshift for clients who are already on the AWS ecosystem and need a managed, petabyte-scale analytical database for business intelligence and reporting.

Looker Studio and Metabase Business intelligence and data visualisation platforms. We use these to build dashboards, reports, and data exploration tools that put analytical insight directly in the hands of the people who need to act on it — without requiring SQL knowledge.

Pandas and NumPy The foundational Python libraries for data manipulation and numerical computation. Every data pipeline, data preprocessing step, and analytical calculation we write in Python starts with these.

Airflow A workflow orchestration platform for scheduling, monitoring, and managing data pipelines. We use Airflow to ensure that data jobs run reliably, in the right order, and with proper error handling and alerting.

Databases

PostgreSQL Our default relational database for most projects. PostgreSQL is open-source, ACID-compliant, and feature-rich — supporting advanced data types, full-text search, JSON storage, and geospatial queries (via the PostGIS extension). It is reliable, well-documented, and has a strong operational ecosystem.

MySQL A widely-used open-source relational database. We use MySQL for projects where existing infrastructure, team familiarity, or third-party integrations make it the right choice. It is fast, stable, and

broadly supported.

MongoDB A document-oriented NoSQL database. MongoDB stores data as JSON-like documents, making it well-suited for applications with flexible or evolving data schemas, high write throughput requirements, or hierarchical data structures. We use it for content management, product catalogues, event logs, and chat message storage.

Redis An in-memory data structure store used as a cache, session store, message broker, and pub/sub engine. Redis is extraordinarily fast — sub-millisecond response times — making it the right tool for rate limiting, session management, real-time leaderboards, job queues, and any data that needs to be accessed repeatedly at high speed.

Elasticsearch A distributed search and analytics engine. We use Elasticsearch for full-text search, faceted filtering, log analysis, and real-time analytics. When a user needs to search across thousands of records with complex filters and ranking — Elasticsearch is what powers that experience.

Supabase An open-source Firebase alternative built on PostgreSQL. Supabase provides a managed database, authentication, real-time subscriptions, storage, and edge functions — all accessible via a clean API. It significantly reduces the time to production for applications that need these capabilities without wanting to configure each service separately.

Firestore (Firestore and Realtime Database) Google's application development platform. Firestore is our choice for real-time, mobile-first applications that need offline support, live synchronisation, and seamless integration with Firebase's authentication and cloud functions. The Realtime Database is used for simpler, lower-latency real-time data scenarios.

Pinecone and Weaviate Vector databases designed for storing and querying high-dimensional embeddings — the numerical representations that AI models use to understand meaning. We use these as the retrieval layer in RAG systems, semantic search, and recommendation engines.

Cloud and Infrastructure

Amazon Web Services (AWS) Our primary cloud platform. We use the full breadth of AWS services including EC2 (virtual machines), ECS and EKS (container orchestration), Lambda (serverless functions), S3 (object storage), RDS (managed databases), CloudFront (CDN), Route 53 (DNS), IAM (identity and access management), SQS (message queuing), SNS (notifications), and many more. We are comfortable architecting, deploying, and managing complex AWS environments — both greenfield and migrated workloads.

Google Cloud Platform (GCP) Used for projects where Google's specific strengths are most relevant — BigQuery for analytics, Vertex AI for ML, Cloud Run for containerised applications, and Firebase for mobile backends. GCP is also the right choice for clients already invested in Google Workspace and Google's data ecosystem.

Microsoft Azure Used for enterprise clients where Microsoft is the strategic technology vendor — particularly for integrations with Microsoft 365, Active Directory, Azure DevOps, and Azure AI services. Azure's compliance certifications also make it the right choice for regulated industries in certain geographies.

Vercel The deployment platform built by the creators of Next.js. Vercel provides zero-configuration deployments, automatic preview environments for every pull request, edge network delivery, and serverless functions — making it the fastest path to production for Next.js applications.

Cloudflare A global content delivery network, DNS provider, DDoS protection service, and edge computing platform. We use Cloudflare in front of every production deployment — it provides automatic

HTTPS, web application firewall (WAF) rules, bot protection, and edge caching that dramatically improves performance and security.

DevOps and CI/CD

Docker A containerisation platform that packages an application and all its dependencies into a portable, reproducible container. Docker ensures that code runs the same way in development, testing, and production — eliminating environment-specific bugs and making deployments predictable and repeatable.

Kubernetes (K8s) The industry-standard container orchestration platform. Kubernetes manages the deployment, scaling, and operation of containerised applications across a cluster of machines. We use Kubernetes for complex, multi-service applications that need fine-grained control over resource allocation, rolling deployments, and auto-scaling.

GitHub Actions Our primary CI/CD automation platform. GitHub Actions runs automated workflows on every code push — running tests, linting code, building container images, and deploying to staging or production environments. Every codebase we deliver is set up with a working CI/CD pipeline on day one.

Terraform An infrastructure-as-code tool that lets us define cloud infrastructure in declarative configuration files. With Terraform, every piece of infrastructure — servers, databases, networking, IAM policies — is version-controlled, peer-reviewed, and reproducible. We never click-to-configure production infrastructure.

Nginx A high-performance web server and reverse proxy. We use Nginx as the entry point for backend applications — handling SSL termination, load balancing, rate limiting, and request routing. It is fast, battle-tested, and highly configurable.

Datadog A cloud monitoring and observability platform. We use Datadog for infrastructure monitoring, application performance monitoring (APM), log management, and alerting. Datadog gives us end-to-end visibility into every layer of a deployed application — from the host machine to the individual database query.

Sentry An error tracking and performance monitoring platform for applications. Sentry captures exceptions in production, provides the full stack trace, the user action that triggered the error, and the environment context — so that bugs can be identified and fixed before users report them.

Prometheus and Grafana Open-source monitoring and visualisation tools. We use Prometheus to collect time-series metrics from application and infrastructure components, and Grafana to build dashboards that make those metrics visible and actionable.

CMS and Commerce Platforms

Shopify The leading e-commerce platform. We build custom Shopify storefronts, themes, and applications — including Shopify Plus implementations for high-volume merchants. We use Shopify's Hydrogen framework for headless commerce builds that need custom frontend experiences while retaining Shopify's commerce engine.

WordPress The world's most widely-used content management system. We build custom WordPress themes, plugins, and WooCommerce integrations for clients who need a flexible, maintainable CMS with a large support ecosystem. We also migrate legacy WordPress sites to modern headless architectures when performance and flexibility demand it.

Payload CMS A modern, TypeScript-native headless CMS built for developers. Payload CMS gives us complete control over the content model, access control, and API — with a clean admin interface that non-

technical teams can use comfortably. We use Payload CMS as our default headless CMS recommendation for new projects.

Strapi An open-source headless CMS built on Node.js. Strapi provides a flexible content model, REST and GraphQL APIs, and a customisable admin panel. It is well-suited for projects that need a self-hosted CMS with strong customisation capability.

Sanity A structured content platform with a real-time content API and a highly customisable editing environment. Sanity is particularly well-suited for content-heavy products where multiple teams are editing simultaneously and the content model needs to be precise and version-controlled.

Contentful A fully managed headless CMS used for enterprise content management. Contentful is the right choice for global organisations that need multi-environment content workflows, localisation support, and enterprise-grade reliability without the overhead of self-hosting.

API and Integration

REST APIs Our default integration approach. We design RESTful APIs that are versioned, documented, and consistent — following established standards for resource modelling, authentication, error handling, and pagination.

GraphQL Used when clients need flexible, efficient data access — particularly in applications with multiple client types consuming data in different shapes and sizes.

gRPC A high-performance remote procedure call framework developed by Google. We use gRPC for internal service-to-service communication in microservices architectures where the performance overhead of HTTP/REST is a constraint.

Webhooks Used for event-driven integrations between systems. We implement webhook producers and consumers for real-time event notification — payment confirmations, status changes, third-party platform events.

Payment Gateways — Razorpay, Stripe Razorpay for Indian market payment processing — supporting UPI, cards, net banking, wallets, and EMI. Stripe for international payment processing with multi-currency support, subscription billing, and marketplace payment splits. We integrate both at the API level, not via plugins — giving full control over the payment experience.

Communication APIs — Twilio, SendGrid Twilio for SMS, voice calls, WhatsApp messaging, and phone number verification. SendGrid for transactional and marketing email delivery at scale. Both integrated at the API level with proper retry logic, delivery tracking, and template management.

Maps and Geolocation — Google Maps API Geolocation, address autocomplete, geocoding, route calculation, and map rendering. Used in logistics, real estate, food delivery, and any application where physical location is a core data dimension.

Identity Verification — DigiLocker, Aadhaar API Indian government-backed identity and document verification. Used for KYC flows in fintech, healthcare, and any regulated application that requires verified identity documents.

Security

JWT (JSON Web Tokens) Stateless authentication tokens used to securely transmit identity and permission claims between services. We use short-lived access tokens with secure refresh token rotation as our standard authentication pattern.

OAuth 2.0 and OpenID Connect Industry-standard protocols for delegated authorisation and identity federation. We implement OAuth 2.0 for third-party integrations and OpenID Connect for single sign-on

(SSO) implementations.

bcrypt and Argon2 Password hashing algorithms. We never store passwords in plain text or with reversible encryption. bcrypt and Argon2 are computationally expensive by design — making brute-force attacks impractical.

AWS IAM and Secrets Manager Identity and access management for cloud resources, and secure storage for credentials, API keys, and configuration secrets. We follow the principle of least privilege — every service gets only the permissions it needs, and credentials are never stored in code or environment files.

SSL/TLS Encryption in transit for all network communication. HTTPS is mandatory on every production deployment. We enforce TLS 1.3 where supported and use HSTS headers to prevent protocol downgrade attacks.

OWASP standards We apply OWASP (Open Web Application Security Project) guidelines across all web and API development — protecting against SQL injection, XSS, CSRF, broken authentication, insecure direct object references, and the other vulnerabilities in the OWASP Top 10.

Testing and Quality Assurance

Jest and Vitest JavaScript testing frameworks for unit and integration tests. We write tests that cover business logic, edge cases, and error paths — not just happy-path coverage numbers.

React Testing Library A testing utility for React components that encourages testing from the user's perspective — interacting with the UI the way a real user would, rather than testing implementation details.

Playwright and Cypress End-to-end testing frameworks that automate browser interaction. We use these to test critical user flows — signup, login, checkout, data entry — across browsers and screen sizes before every deployment.

Pytest Python's standard testing framework. Used for testing FastAPI and Django applications, ML model behaviour, and data pipeline logic.

Postman and Insomnia API testing and documentation tools. Every API we build is documented and tested through these tools — covering success responses, error responses, authentication, and edge cases.

Design and Collaboration

Figma Our primary design and prototyping tool. Every UI project starts in Figma — from wireframes and user flows to high-fidelity designs and interactive prototypes. We deliver complete design systems in Figma with reusable components, design tokens, and detailed specifications for development handoff.

Adobe Creative Suite Used for brand identity work, illustration, motion graphics, and assets that require capabilities beyond UI design — photography editing, vector illustration, and video production.

Whimsical Used for flowcharts, wireframes, mind maps, and system diagrams during the discovery and architecture phase of projects.

Notion and Linear Notion for documentation, knowledge bases, and project wikis. Linear for engineering project management — issues, sprints, milestones, and roadmaps.

Slack and Loom Slack for asynchronous team and client communication. Loom for recorded video updates, screen walkthroughs, and handoff documentation.

We work with new technologies as they mature and prove themselves in production environments. The stack above reflects what we use today — and it evolves as the landscape evolves. If you are working with a

technology not listed here, ask us. The answer is likely yes.

Our Internal Products

In addition to client work, OrynticLabs builds and maintains its own product suite. These products are used internally, offered to clients as part of engagements, and in some cases available as standalone products.

OryAI

OryAI is OrynticLabs' proprietary AI intelligence platform. It is a modular system for building, deploying, and managing custom AI agents, automation pipelines, and large language model integrations inside a business.

OryAI is the infrastructure layer that sits underneath the AI features we build for clients. It provides the agent framework, the knowledge management system, the evaluation layer, and the integration connectors — so that AI capabilities can be deployed reliably in production, not just demonstrated in a prototype.

What OryAI includes:

- Custom AI agent builder and orchestration engine
- RAG system with support for multiple vector databases
- LLM integration layer (OpenAI, Anthropic, open-source models)
- Workflow automation builder
- Conversation memory and session management
- Evaluation and monitoring dashboard
- Multi-language support including Hindi and English

OryCMS

OryCMS is OrynticLabs' headless content management system, built for teams who need the flexibility of a modern content architecture without the complexity of setting up and maintaining open-source solutions.

OryCMS is designed for product teams and marketing teams that need to publish, manage, and distribute content across multiple channels — web, mobile, email, third-party platforms — without depending on a developer for every change.

What OryCMS includes:

- Headless API-first architecture with REST and GraphQL support
- Visual content editor with structured content modelling
- Role-based access control for editorial workflows
- Multi-site and multi-language content management
- Media library with image optimisation and CDN delivery
- Webhook and integration support for popular platforms
- White-label ready for agency and enterprise deployments

PerformX

PerformX is OrynticLabs' internal Human Resource Management System (HRMS), built to manage the full employee lifecycle — from onboarding through payroll, leave management, performance reviews, and

offboarding.

PerformX was built because the existing HRMS tools available in the market were either too expensive, too rigid, or missing features specific to how OrynticLabs operates. We built the version we wanted to use ourselves — and we now offer it to other companies that have the same frustration.

What PerformX includes:

Employee records and document management

Onboarding workflow automation

Attendance tracking and leave management with automated LOP calculation

Payroll processing with salary structure management and payslip generation

Performance review cycles with rating and goal management

Internal mail and meeting management system

Role-based access for HR, managers, and employees

Payment gateway integration for direct salary disbursement

Deployed on AWS with full data security and backup infrastructure

Resource and Staffing Services

OrynticLabs provides two additional service models for companies that need technology talent without building a full internal team.

Staff Augmentation

We provide experienced engineers, designers, and product managers who embed directly into your team and work under your direction. This is for companies that have a clear product vision but need additional capacity to execute it.

Every person we place from OrynticLabs is a full-time member of our team — vetted, trained in our engineering standards, and backed by our technical leadership. You get the output of a strong hire without the overhead of a permanent headcount.

Available roles:

Frontend and backend engineers

Full-stack developers

Mobile developers

AI and ML engineers

UI/UX designers

QA engineers

DevOps and cloud engineers

Product managers

Technology Consulting

We work with leadership teams — founders, CTOs, heads of product — to help them make better technology decisions. This is advisory work: helping you understand your options, evaluate your current architecture, plan a migration, assess a vendor, or define a technology roadmap.

Consulting engagements are structured, time-bound, and outcome-focused. We do not bill indefinitely for advice. We come in, understand the situation, provide a clear recommendation with our reasoning, and leave you with something actionable.

How We Work

Every OrynticLabs engagement — whether it is a small web build or a multi-year product partnership — follows the same foundational approach.

Understand before recommending. We spend time understanding your business, your users, your constraints, and your goals before we suggest a solution. The wrong solution delivered well is still the wrong solution.

Scope before building. Every engagement begins with a documented scope of work. You know exactly what you are getting, what it will cost, and what success looks like — before any work begins.

Show progress every week. We work in two-week sprints with a working demo every Friday. You see real progress every week, not a finished product at the end of six months.

Be honest when something is not working. If a technical approach is not the right one, we say so — even if we already started down that path. If timelines need to change, we communicate early, not when it is too late to adjust.

Own the outcome, not just the output. We measure our success by whether the product works for your users and moves your business forward. Shipping code is not the goal. The goal is the result that code creates.

Why OrynticLabs

There are a lot of technology companies. Here is what is genuinely different about working with us.

We cover the full stack. Web, mobile, AI, data, design, infrastructure, CMS, commerce, and consulting — under one roof, by one accountable team. You do not need to coordinate between a design agency, a development agency, and an AI consultant who have never spoken to each other.

We build for the long term. We write code that the next engineer can understand and extend. We design systems that can scale. We document what we build. We do not optimise for the speed of the first delivery at the expense of everything that comes after it.

We have real AI capability. Not prompt engineering. Not ChatGPT wrappers. We build AI systems — trained models, intelligent agents, data pipelines, and production-grade ML infrastructure. If you want genuine AI in your product, this is who you call.

We treat your business like our own. The people on your kickoff call are the people building your product. Our founders are directly involved in the work. Your project does not get handed off to a junior team once the contract is signed.

We have internal products. OryAI, OryCMS, and PerformX are not theoretical — they are systems we use ourselves. That means we understand the operational reality of what we build, not just the technical specification.

Contact

OrynticLabs Private Limited Incorporated in India under the Companies Act, 2013

Website: www.orynticlabs.com Sales: sales@orynticlabs.com Support: support@orynticlabs.com General: hello@orynticlabs.com Phone: +91 79017 17617

This document is the property of OrynticLabs Private Limited. It is intended for authorised recipients only and may not be reproduced, distributed, or shared without prior written permission.

© 2026 OrynticLabs Private Limited. All rights reserved.